Extra Credit Assignment

The code for the assignment is run using NYU CIMS's crackle1 server which has the Intel Xeon E5630 CPU.
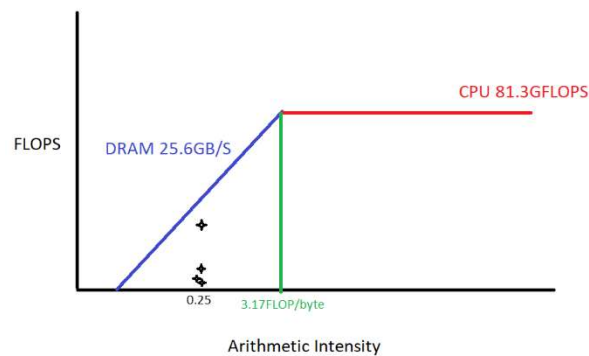

Figure: roofline model for Q1, Q2, Q3

Q1. The result for the unrolled loop dot product is:

time:1.262464secs bandwidth:6.339985GB/s flops:1.584204GFLOP/s arithmetic_intensity:0.249875FLOP/byte

Q2. The L1 and L2 caches are 128KB and 1024KB, respectively (found from this link). To get our blocks to fit within the caches we must first find the appropriate size of the blocks. Since we will be using single-point precision floats (4 bytes) and are storing 2 submatrices in "shared" memory (i.e. cache in this case), we have $sqrt\left(\frac{128,000byte}{2*4byte}\right) = 126.491 => 100$x100 Block dimension for L1-cache. Similarly, for L2-cache, we get 250x250. The result for the tiling matric multiplication routine for BlockDim = 100x100 is:

time:2.042972secs bandwidth:3.956981GB/s flops:0.983861GFLOP/s arithmetic_intensity:0.248639FLOP/byte

For 250x250:

time:1.932407secs bandwidth:4.158543GB/s flops:1.037048GFLOP/s arithmetic_intensity:0.249378FLOP/byte

Q3. I wasn't able to find the actual implementation of cblas_sgemm other than that it performs $O(n^3)$. Therefore, I assume that it uses the standard matrix multiply procedure but with hardware optimizations. The result for the MKL matrix multiply routine is:

time:0.154507secs bandwidth:51.803576GB/s flops:12.944422GFLOP/s arithmetic_intensity:0.249875FLOP/byte

Notes:

It is interesting to note that our problem size is 1000x1000x4bytes = 4MB, which happens to fit inside the L3 cache of 12MB. So, the roofline model (which I've drawn from the lecture notes since I am using the same CPU) might be inaccurate for our example. L3 cache is faster than DRAM, the scale of which I have not been able to determine from Intel's spec sheet or anywhere else.